

# Real-Time Robot Self-Modeling via Direct Sensorimotor Decoding

*Overcoming the Neural Rendering Latency Bottleneck*

Muhammad Zeeshan Asghar

Master's Program in Data Science  
Higher School of Economics

December 17, 2025

**Course:** Machine Learning and Data Mining Implementation Project

---

# The Reality Gap: Why Self-Modeling?

- ① **The Problem:** Robots in unstructured environments (Space, Rescue) suffer damage.

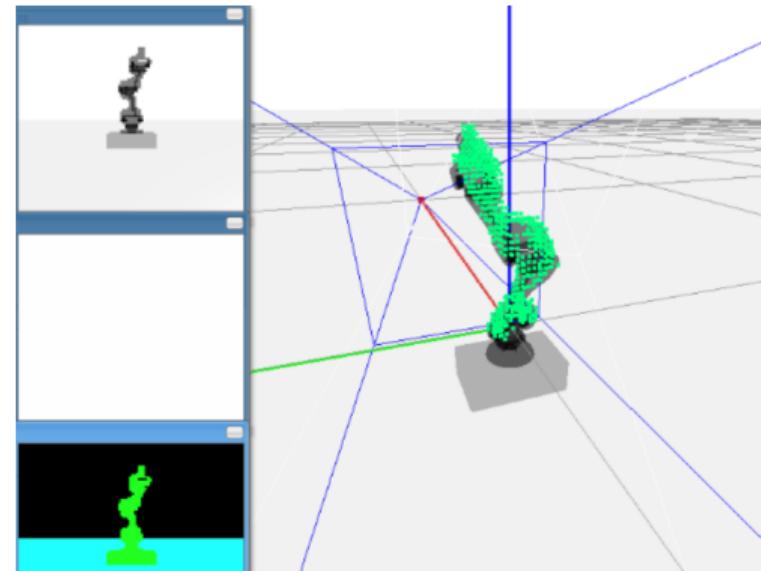


Figure 1: 4-DOF Manipulator in PyBullet Simulation Environment.

# The Reality Gap: Why Self-Modeling?

- ① **The Problem:** Robots in unstructured environments (Space, Rescue) suffer damage.
- ② **The Failure Mode:** Traditional controllers use fixed kinematic files (URDF). If the body bends, the model drifts → **Catastrophic Failure**.

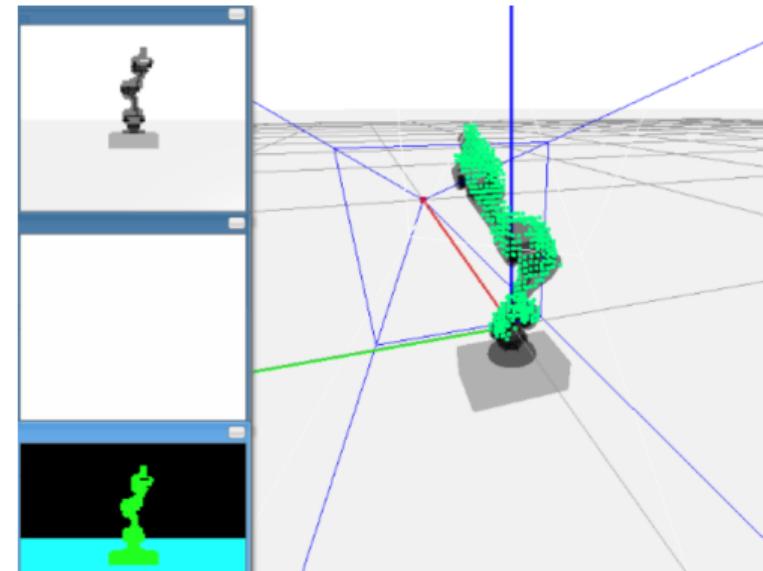


Figure 1: 4-DOF Manipulator in PyBullet Simulation Environment.

# The Reality Gap: Why Self-Modeling?

- ➊ **The Problem:** Robots in unstructured environments (Space, Rescue) suffer damage.
- ➋ **The Failure Mode:** Traditional controllers use fixed kinematic files (URDF). If the body bends, the model drifts → **Catastrophic Failure**.
- ➌ **The Goal:** A robot that learns its own shape from vision in real-time.

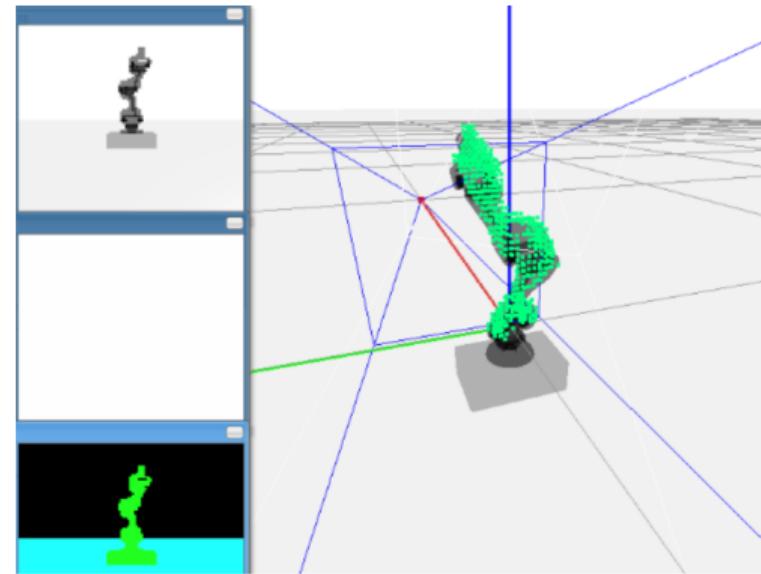


Figure 1: 4-DOF Manipulator in PyBullet Simulation Environment.

# The Reality Gap: Why Self-Modeling?

- ① **The Problem:** Robots in unstructured environments (Space, Rescue) suffer damage.
- ② **The Failure Mode:** Traditional controllers use fixed kinematic files (URDF). If the body bends, the model drifts → **Catastrophic Failure**.
- ③ **The Goal:** A robot that learns its own shape from vision in real-time.
- ④ **The Bottleneck:**

## Current SOTA - FFKSM Nature 2024

Uses Neural Radiance Fields (NeRF).

**Latency:** ~5.22 FPS (192ms).

**Required:** >1000 FPS (1ms) for control loops.

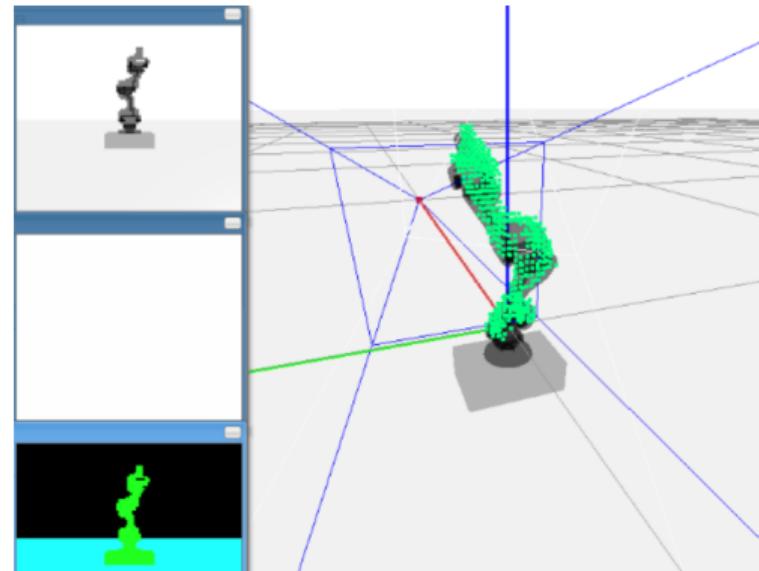


Figure 1: 4-DOF Manipulator in PyBullet Simulation Environment.

# Project Contributions & Scope

## 1. Reproduction

### FFKSM (Baseline)

Implemented from scratch.  
Solved "Kinematic Blindness" Class  
Imbalance bugs.

**Result: 5.22 FPS**

## 2. Failed Hypothesis

### K-3DGS (Explicit)

Tried Kinematic 3D Gaussians.  
Failed anisotropic optimization.

**Result: "Blobby" Artifacts**

## 3. Innovation

### NeuroKin (Ours)

Direct Sensorimotor Decoding.  
Bypassed 3D reconstruction entirely.

**Result: 7,400 FPS**

# Data Generation via Chaotic Dynamics

**Problem:** Random motor babbling produces jerky, discontinuous motion that fails to capture kinematic dependencies.

**Solution:** We employ **Lorenz Attractors** to generate smooth, ergodic trajectories that efficiently explore the workspace.

$$\begin{aligned}\dot{x} &= \sigma(y - x) \\ \dot{y} &= x(\rho - z) - y \\ \dot{z} &= xy - \beta z\end{aligned}$$

Parameters:  $\sigma = 10, \rho = 28, \beta = 8/3$ .

Generates deterministic but non-repeating joint angles for 4-DOF.

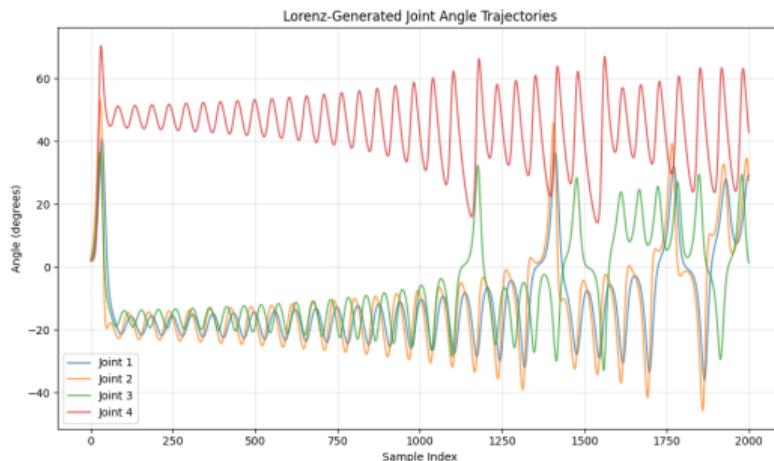


Figure 2: Chaotic trajectories (2,000 samples) ensuring ergodic workspace coverage.

# Data Processing Pipeline

We constructed a robust pipeline to convert simulation output into training tensors.

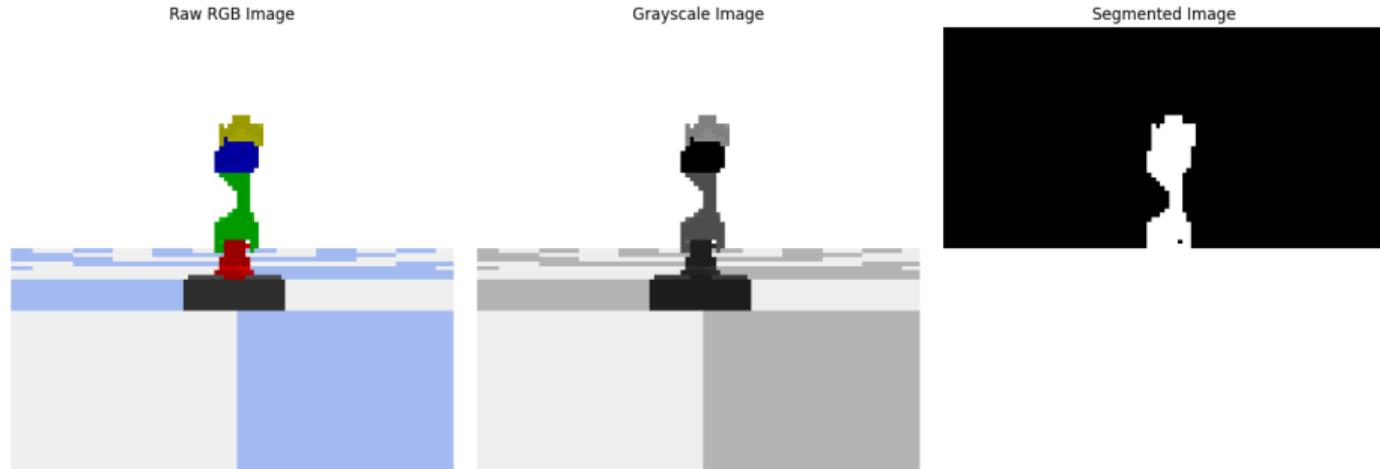


Figure 3: Data preprocessing pipeline converting RGB input to binary training masks.

| Input                                  | Processing                          | Output                            |
|--|-------------------------------------|-----------------------------------|
| PyBullet RGB ( $H \times W \times 3$ ) | Thresholding ( $>240$ ) & Grayscale | Binary Masks ( $100 \times 100$ ) |

Dataset Scale: 2,000 samples (1,600 Train / 400 Val). This constrained regime tests data efficiency.

# Method 1: FFKSM Architecture (Reproduction)

**Approach:** Volumetric Querying. Condition Density  $\sigma$  on Joints  $\theta$ .

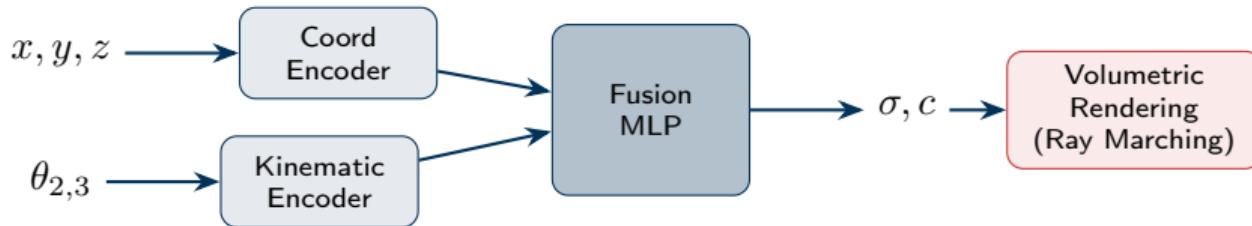


Figure 4: FFKSM architecture with Split-Encoders and Volumetric Rendering head.

- **Split Encoder:** Decouples kinematics from geometry.
- **Bottleneck:** Ray Marching = 640k evals/image.

## Results

**PSNR:** 17.35 dB

**Speed:** 5.22 FPS (Too slow)

# Engineering Challenge: The "Black Screen" Convergence

## The Pathology:

- Robot occupies only  $\sim 15\%$  of pixels.
- Network converged to **all-zeros** (local minima, Loss  $\approx 0.15$ ).

## The Solution: Curriculum Learning

- **Phase 1** ( $t < 500$ ): Train ONLY on center  $50 \times 50$  crop (Occupancy  $> 40\%$ ).
- **Phase 2** ( $t \geq 500$ ): Expand to full image.

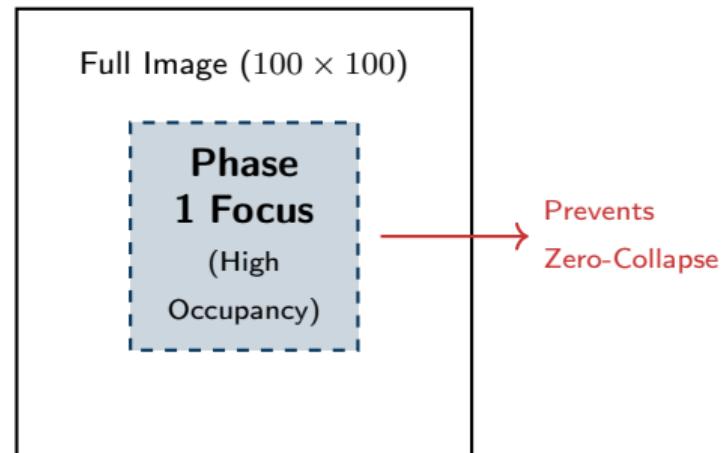


Figure 5: Curriculum Learning Strategy.

# Method 2: K-3DGS Architecture (Failed Hypothesis)

**Hypothesis:** Rasterization is faster than Ray-Marching.

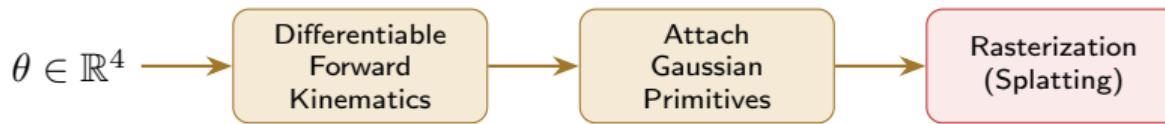


Figure 6: Kinematic 3D Gaussian Splatting (K-3DGS) pipeline.

## Optimization Failure

**Anisotropic:** Too slow and unstable.

**Isotropic:** Stable, but "blobby".

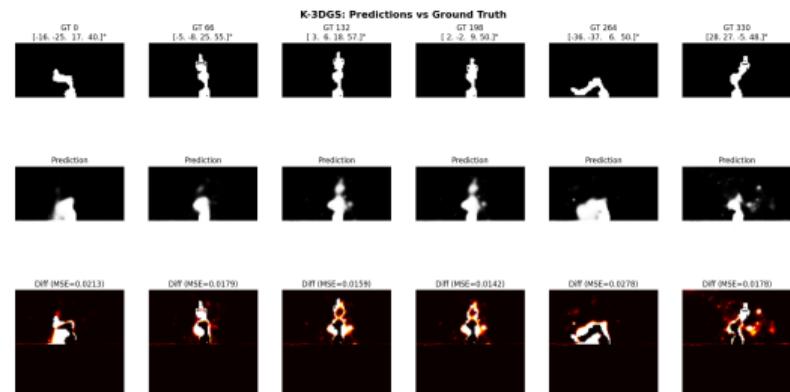


Figure 7: "Blobby" artifacts due to isotropic constraints (17.01 dB).

# Method 3: NeuroKin Architecture (Success)

**The Pivot:** Bypassing 3D reconstruction for Direct Sensorimotor Decoding.

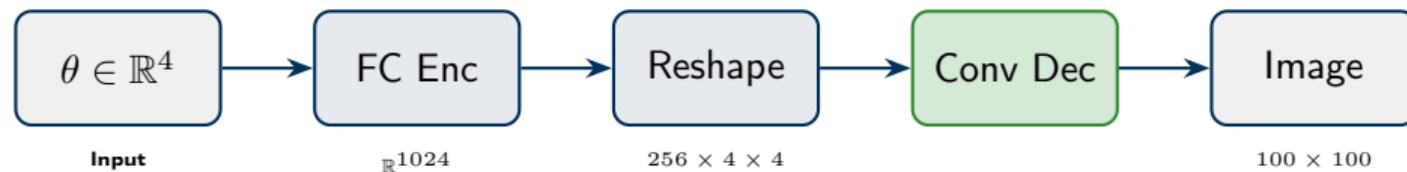


Figure 8: NeuroKin Direct Decoding Architecture.

Why it wins:

- **O(1) Complexity:** Single pass.
- **Dense Supervision:** 10,000 pixels updated per eval.

## Performance

**Speed:** 7,400 FPS (1400x faster)  
**Quality:** 21.88 dB (+4.53 dB)

# Method 4: ResNeuroKin-D (Multi-Task Extension)

**Motivation:** Pure NeuroKin lacks geometric structure in the latent space. **Solution:** Force the network to learn *Synthetic Heuristic Depth* alongside silhouettes.

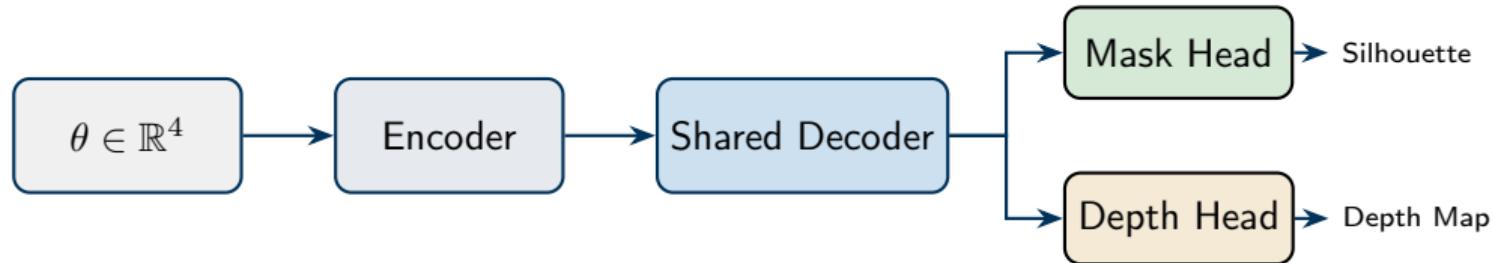


Figure 9: ResNeuroKin-D Dual-Head Architecture.

- **Geometric Prior:** Depth loss forces spatial awareness.
- **Trade-off:** 2,500 FPS (Slower than NeuroKin, but still > 1 kHz).
- **Heuristic Depth:** Generated via kinematic distance.

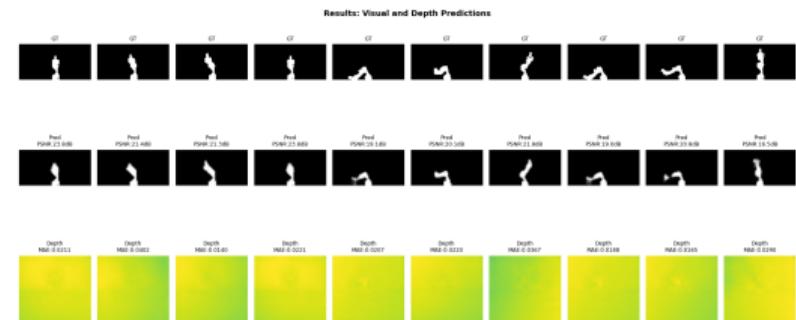


Figure 10: Dual output predictions (Silhouette + Depth).

# Quantitative Comparison of Self-Modeling Approaches

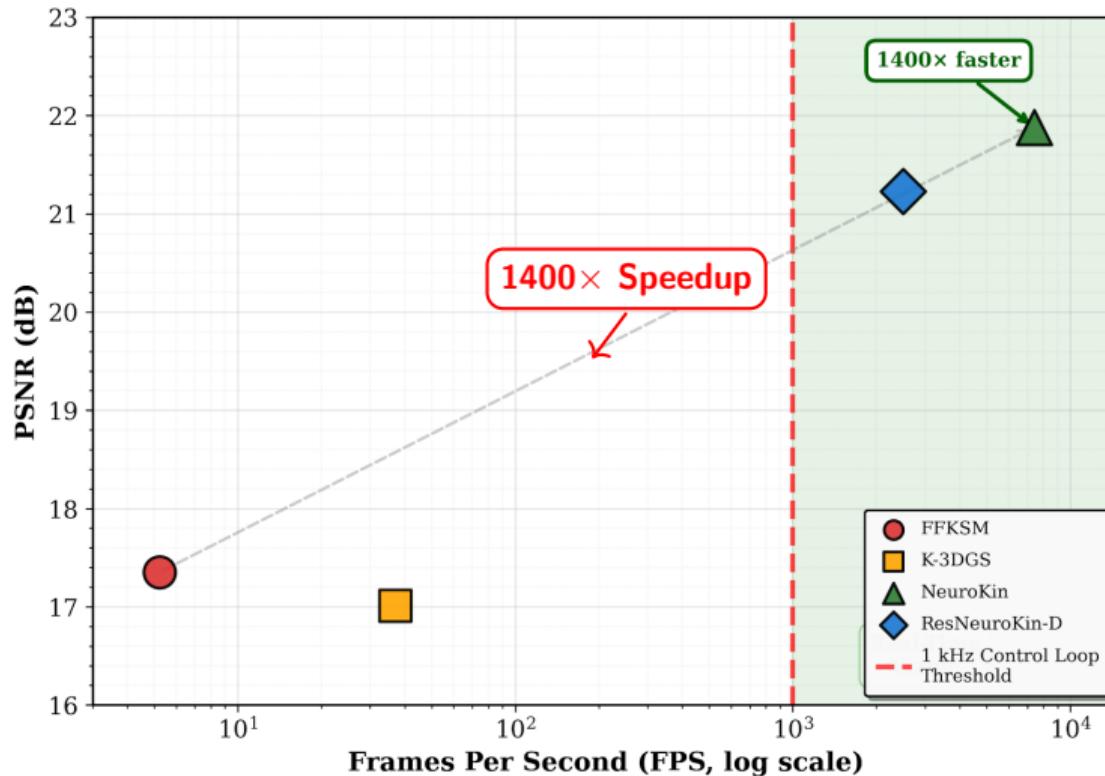
Table 1: Quantitative comparison of self-modeling architectures.

| Method                 | PSNR            | FPS         | Latency         | Train Time     | Speedup      |
|------------------------|-----------------|-------------|-----------------|----------------|--------------|
| FFKSM                  | 17.35 dB        | 5.22        | 191.6 ms        | 50.0 min       | 1.0×         |
| K-3DGS                 | 17.01 dB        | 37          | 27.0 ms         | 0.8 min        | 7.1×         |
| <b>NeuroKin (Ours)</b> | <b>21.88 dB</b> | <b>7400</b> | <b>0.135 ms</b> | <b>0.5 min</b> | <b>1418×</b> |
| ResNeuroKin-D          | 21.23 dB        | 2500        | 0.4 ms          | 6.0 min        | 479×         |

- **NeuroKin** achieves a massive **1418x speedup** over the baseline.
- **FFKSM** latency (192ms) is unusable for real-time control.
- **ResNeuroKin-D** maintains >400x speedup while learning depth.

# The Pareto Frontier: Speed vs Quality

Speed-Quality Pareto Frontier: Self-Modelling Approaches



## Key Result

- NeuroKin** dominates:
- 7,400 FPS
  - +4.53 dB PSNR

# Ablation: Data Efficiency (RQ3)

**Question:** How much data does NeuroKin actually need? **Answer:** NeuroKin beats FFKSM with only **25%** of the data.

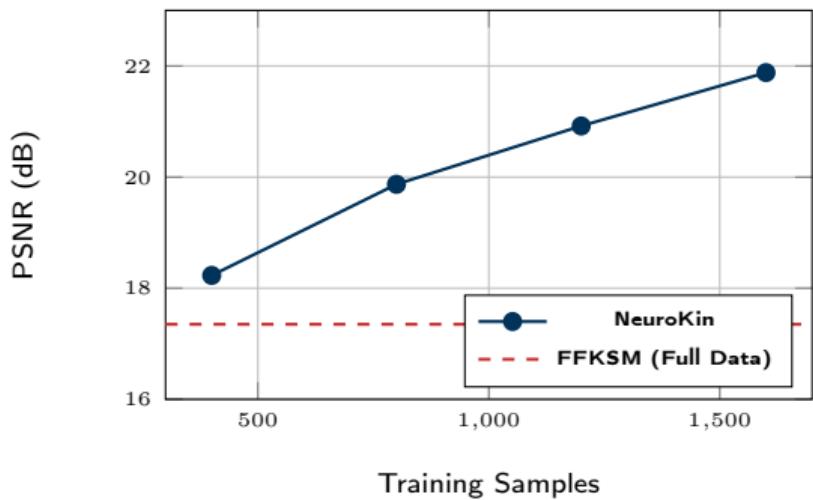


Figure 12: Performance vs Dataset Size.

## Key Insight

With just **400 samples**, NeuroKin achieves **18.23 dB**, exceeding FFKSM's performance on the full 1,600 dataset.

This confirms that direct decoding is significantly more **sample efficient** than volumetric rendering.

# Qualitative Comparison

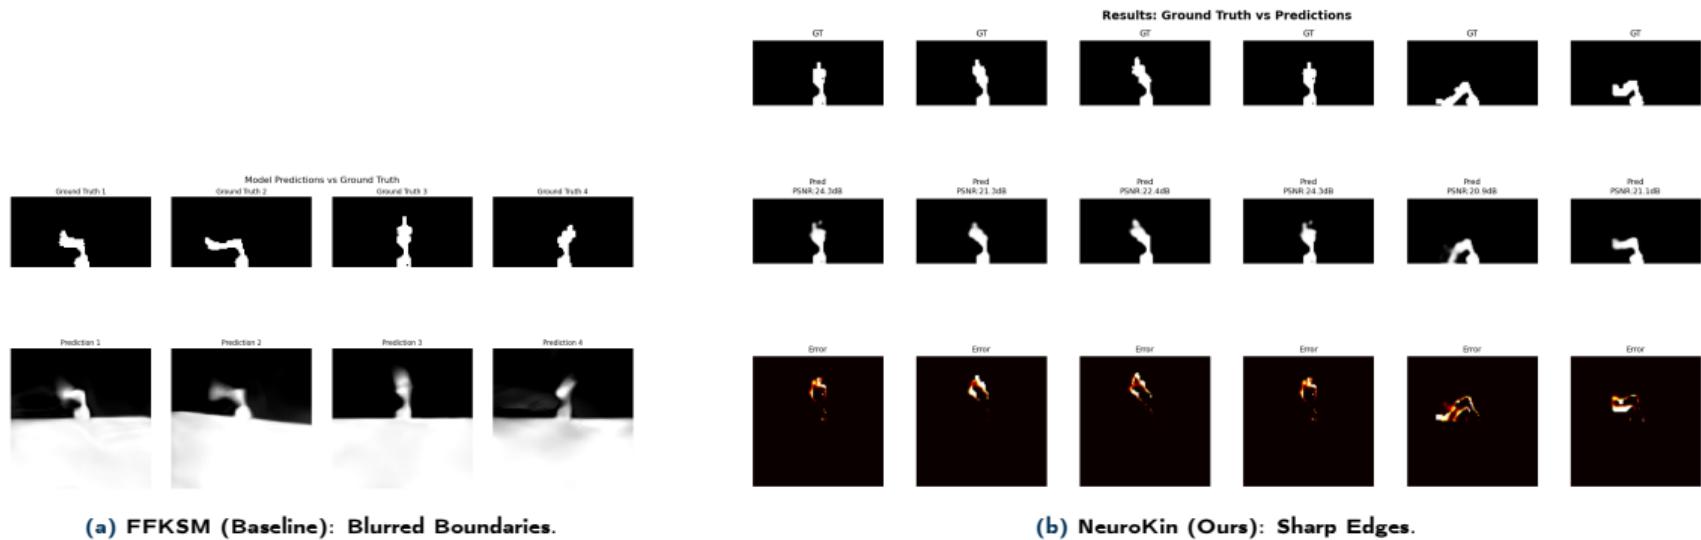


Figure 13: Qualitative comparison of generated robot silhouettes.

# Future Directions: Towards Resilient Autonomy

- **Multi-View Generalization:**

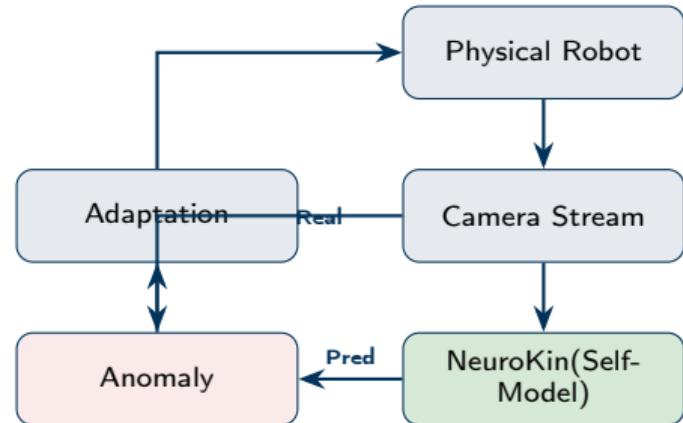
- ▶ Current limitation: Single fixed camera.
- ▶ Proposal: Train on **multi-view datasets** to learn a view-invariant latent representation.

- **Evolving Sparse Topologies:**

- ▶ Use **Evolutionary Algorithms (NEAT)** to prune the NeuroKin network.
- ▶ Goal: Find the "minimal viable brain" for microcontrollers.

- **Real-Time Damage Adaptation Loop:**

- ▶ Close the control loop: Use the <1ms inference to detect damage and adapt online.



# Conclusion & Impact

- ❶ **Rigorous Reproduction:** Validated FFKSM, exposed latency floor (5.22 FPS).
- ❷ **Negative Result:** Kinematic 3DGS is unstable without anisotropic constraints.
- ❸ **SOTA Performance:** NeuroKin (7,400 FPS).

## Why it matters?

By discarding 3D reconstruction, we enable **1 kHz control loops** on edge hardware (Jetson Orin), allowing robots to detect damage in < 1 ms.

Code available: [https://github.com/YoloPopo/robot\\_self\\_modelling](https://github.com/YoloPopo/robot_self_modelling)

# Thank You!

Questions?

*Muhammad Zeeshan Asghar  
Higher School of Economics*